

Stats 102C - Lecture 4-2: Box-Mueller Sampling and Multivariate Sampling

Michael Tsiang (adapted by Miles Chen)

Department of Statistics

Week 4 Wednesday

UCLA

Section 1

Box-Muller Transform

Box-Muller Transform for Normal Distribution

- The Box-Muller transform (George Box and Mervin Muller, 1958) is a method to transform two uniform random variables into a pair of independent standard normal random variables.
- The main idea is to change coordinates from Cartesian to polar coordinates.

Box-Muller Transform

Let X and Y be independent standard normal random variables.

The joint PDF of X and Y is given by

$$\begin{aligned}f_{xy}(x, y) &= f(x)f(y) \\&= \frac{1}{\sqrt{2\pi}}e^{-x^2/2} \cdot \frac{1}{\sqrt{2\pi}}e^{-y^2/2} \\&= \frac{1}{2\pi}e^{-\frac{x^2+y^2}{2}}\end{aligned}$$

Box-Muller Transform

The relationship between Cartesian coordinates (x, y) and polar coordinates (r, θ)

$$x = r \cos \theta$$

$$y = r \sin \theta$$

We change $f_{xy}(x, y)$ to polar coordinates:

$$f_{XY}(x, y) dx dy = f_{R\Theta}(r, \theta) dr d\theta$$

Box-Muller Transform

It's probably been a while, but from multivariate calculus, the change of variables is given by:

$$\begin{aligned}f_{R\Theta}(r, \theta) &= f_{XY}(x, y) \frac{dxdy}{drd\theta} \\&= f_{XY}(x, y) \underbrace{\left| \frac{\partial(x, y)}{\partial(r, \theta)} \right|}_{= J},\end{aligned}$$

where J is the Jacobian

$$J = \begin{vmatrix} \frac{\partial x}{\partial r} & \frac{\partial x}{\partial \theta} \\ \frac{\partial y}{\partial r} & \frac{\partial y}{\partial \theta} \end{vmatrix} = \begin{vmatrix} \cos \theta & -r \sin \theta \\ \sin \theta & r \cos \theta \end{vmatrix} = r \cos^2 \theta + r \sin^2 \theta = r$$

Since $x^2 + y^2 = r^2$, then

$$f_{R\Theta}(r, \theta) = f_{XY}(x, y) \cdot r = \frac{1}{2\pi} e^{-\frac{r^2}{2}} \cdot r$$

For $r \geq 0$ and $\theta \in [0, 2\pi)$, we have

$$f_{R\Theta}(r, \theta) dr d\theta = \frac{1}{2\pi} e^{-\frac{r^2}{2}} \cdot r dr d\theta$$

Applying another change of variables from (r, θ) to (r^2, θ) , we have

$$dr^2 = 2rdr, \text{ so } rdr = \frac{1}{2} dr^2$$

The joint PDF $f_{R\Theta}(r, \theta)drd\theta$, for $r \geq 0$ and $\theta \in [0, 2\pi)$, can now be written as

$$\begin{aligned} f_{R\Theta}(r, \theta)drd\theta &= f_{R^2\Theta}(r^2, \theta)dr^2d\theta \\ \frac{1}{2\pi}e^{-\frac{r^2}{2}} \cdot r dr d\theta &= \frac{1}{2\pi}e^{-\frac{r^2}{2}} \cdot \frac{1}{2}dr^2d\theta \\ &= \left(\frac{1}{2}e^{-\frac{r^2}{2}}dr^2\right)\left(\frac{1}{2\pi}d\theta\right) \\ &= f_{R^2}(r^2)dr^2 \cdot f_\Theta(\theta)d\theta, \end{aligned}$$

which shows:

- $R^2 \perp \Theta$ (i.e., R^2 and Θ are independent)
- $\Theta \sim \text{Unif}(0, 2\pi)$
- $R^2 \sim \text{Exp}(\lambda = \frac{1}{2})$

- Generate $\Theta \sim \text{Unif}(0, 2\pi)$
- Generate $V \sim \text{Exp}(\lambda = \frac{1}{2})$ (i.e., $V = R^2$) and compute

$$R = \sqrt{V}$$

- Compute

$$X = R \cos \Theta$$

$$Y = R \sin \Theta$$

Then $X, Y \sim \mathcal{N}(0, 1)$

Note that we can use $U \sim \text{Unif}(0, 1)$ to sample from $\text{Exp}(\lambda = \frac{1}{2})$ using the inverse CDF method:

$$-\frac{1}{\lambda} \log U = -2 \log U \sim \text{Exp}(\lambda = \frac{1}{2})$$

Box-Muller algorithm:

- ① Generate $U \sim \text{Unif}(0, 1)$ and compute

$$\Theta = 2\pi U$$

- ② Generate $V \sim \text{Unif}(0, 1)$ and compute

$$R = \sqrt{-2 \log V}$$

- ③ Compute

$$X = R \cos \Theta$$

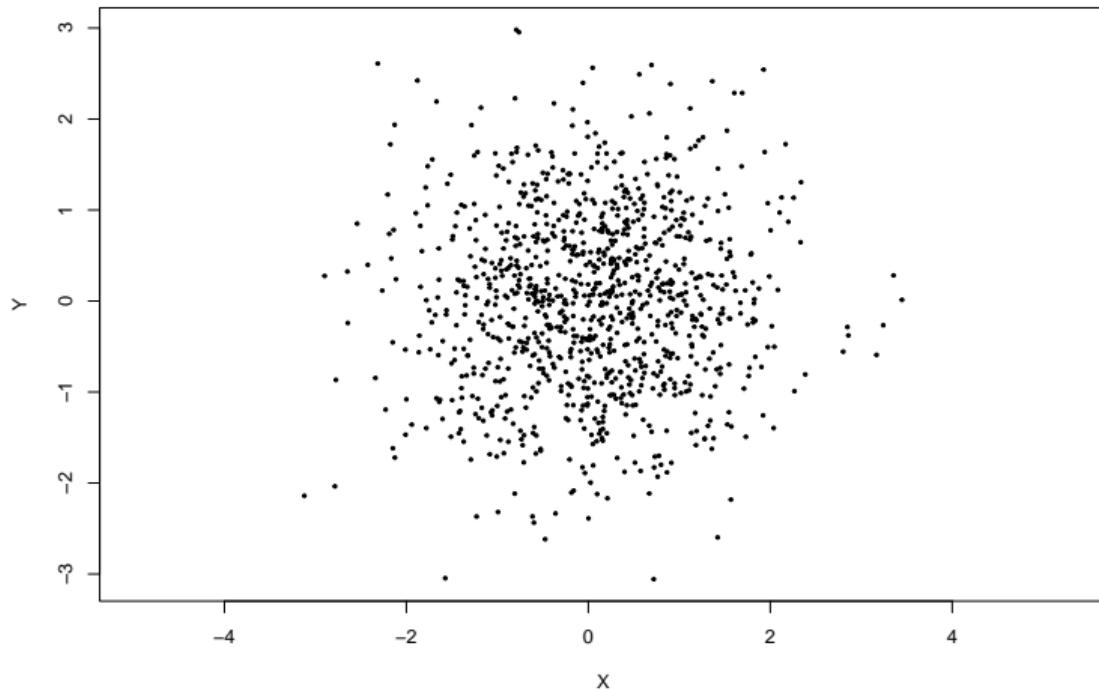
$$Y = R \sin \Theta$$

Then $X, Y \sim \mathcal{N}(0, 1)$

Copyright Miles Chen. For personal use only. Do not distribute.

```
set.seed(9999)
n <- 1000
# Generate n points from Unif(0,1)
U <- runif(n,0,1)
# Compute Theta
Theta <- 2*pi*U
# Generate n points from Unif(0,1)
V <- runif(n,0,1)
# Compute R
R <- sqrt(-2*log(V))
# Compute X and Y
X <- R*cos(Theta)
Y <- R*sin(Theta)
```

```
plot(X, Y, cex = 0.4, asp = 1, pch = 19)
```



https://en.wikipedia.org/wiki/Box%E2%80%93Muller_transform

https://upload.wikimedia.org/wikipedia/commons/1/1f/Box-Muller_transform_visualisation.svg

Normal Distribution

We can now assume we can generate samples from $N(0, 1)$.

How can we use samples from $N(0, 1)$ to sample from

$W \sim N(\mu, \sigma^2)$, for any $\mu \in R, \sigma^2 > 0$?

Normal Distribution

- Generate $Z \sim N(0, 1)$.
- Then $W = \mu + \sigma Z \sim N(\mu, \sigma^2)$.

Why?

- W is normally distributed.
- $E(W) = E(\mu + \sigma Z) = \mu + \sigma E(Z) = \mu$
- $Var(W) = Var(\mu + \sigma Z) = Var(\sigma Z) = \sigma^2 Var(Z) = \sigma^2$

Section 2

Generating a Multivariate Sample

Multivariate Sampling

Sometimes we need to generate random values from a joint distribution of multiple random variables.

We'll look at generating data from bivariate distributions.

Eventually, we'll want to generate samples from higher dimensional distributions.

Bivariate Discrete Distribution

Let's say random variables X and Y have the following joint PMF:

X/Y	$y = 0$	$y = 1$
$x = 0$	0.2	0.6
$x = 1$	0.1	0.1

How can we generate samples from this distribution?

Bivariate Discrete Distribution

We map the pairs X and Y onto interval $(0, 1)$.

Let $p_{xy} = \Pr(X = x, Y = y)$

$$F_1 = p_{00} = 0.2$$

$$F_2 = p_{00} + p_{10} = 0.3$$

$$F_3 = p_{00} + p_{10} + p_{01} = 0.9$$

$$F_4 = p_{00} + p_{10} + p_{01} + p_{11} = 1.0$$

① Generate $U \sim \text{Unif}(0, 1)$

② Then

$$(X, Y) = \begin{cases} (0, 0) & \text{if } 0 < U \leq F_1 \\ (1, 0) & \text{if } F_1 < U \leq F_2 \\ (0, 1) & \text{if } F_2 < U \leq F_3 \\ (1, 1) & \text{if } F_3 < U < 1 \end{cases}$$

A more complex distribution

X / Y	y = 1	y = 2	y = 3	y = 4
x = 1	0.02	0.06	0.03	0.05
x = 2	0.03	0.04	0.02	0.05
x = 3	0.04	0.2	0.1	0.1
x = 4	0.01	0.1	0.05	0.1

$$F_1 = p_{11} = 0.02$$

$$F_2 = p_{11} + p_{21} = 0.05$$

$$F_3 = p_{11} + p_{21} + p_{31} = 0.09$$

$$\vdots = \vdots = \vdots$$

$$F_{15} = p_{11} + p_{21} + p_{31} + \dots + p_{34} = 0.9$$

$$F_{16} = p_{11} + p_{21} + p_{31} + \dots + p_{34} + p_{44} = 1.0$$

① Generate $U \sim \text{Unif}(0, 1)$

② Then

$$(X, Y) = \begin{cases} (1, 1) & \text{if } 0 < U \leq F_1 \\ (2, 1) & \text{if } F_1 < U \leq F_2 \\ \vdots & \vdots \\ (3, 4) & \text{if } F_{14} < U \leq F_{15} \\ (4, 4) & \text{if } F_{15} < U < 1 \end{cases}$$

Code Part 1

```
prob_matrix <- matrix(  
  c(0.02, 0.06, 0.03, 0.05,  
    0.03, 0.04, 0.02, 0.05,  
    0.04, 0.2 , 0.1, 0.1,  
    0.01, 0.1 , 0.05, 0.1),  
  byrow = TRUE, nrow = 4)  
  
cum_prob <- cumsum(as.vector(prob_matrix))  
  
cum_prob  
  
## [1] 0.02 0.05 0.09 0.10 0.16 0.20 0.40 0.50 0.53 0.55 0.65 0.70 0.75 0.80 0.90 1.00
```

```
prob_matrix
```

```
##      [,1] [,2] [,3] [,4]
## [1,] 0.02 0.06 0.03 0.05
## [2,] 0.03 0.04 0.02 0.05
## [3,] 0.04 0.20 0.10 0.10
## [4,] 0.01 0.10 0.05 0.10

x_vals <- rep(1:4, 4)
y_vals <- rep(1:4, each = 4)
print(rbind(cum_prob, x_vals, y_vals))

##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14] [,15] [,16]
## cum_prob 0.02 0.05 0.09 0.1 0.16 0.2 0.4 0.5 0.53 0.55 0.65 0.7 0.75 0.8 0.9 1
## x_vals   1.00 2.00 3.00 4.0 1.00 2.0 3.0 4.0 1.00 2.00 3.00 4.0 1.00 2.0 3.0 4
## y_vals   1.00 1.00 1.00 1.0 2.00 2.0 2.0 2.0 3.00 3.00 3.00 3.0 4.00 4.0 4.0 4
```

```

# one value
u <- 0.01
print(rbind(cum_prob, x_vals, y_vals))

## [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14] [,15] [,16]
## cum_prob 0.02 0.05 0.09 0.1 0.16 0.2 0.4 0.5 0.53 0.55 0.65 0.7 0.75 0.8 0.9 1
## x_vals   1.00 2.00 3.00 4.0 1.00 2.0 3.0 4.0 1.00 2.00 3.00 4.0 1.00 2.0 3.0 4
## y_vals   1.00 1.00 1.00 1.0 2.00 2.0 2.0 2.0 3.00 3.00 3.00 3.0 4.00 4.0 4.0 4

position <- min(which( u < cum_prob)) # compare u to cum_prob and finds the first value that is greater
position

## [1] 1
x_vals[position] # value to save for X

## [1] 1
y_vals[position] # value to save for Y

## [1] 1

```

```

# one value
u <- 0.03
print(rbind(cum_prob, x_vals, y_vals))

## [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14] [,15] [,16]
## cum_prob 0.02 0.05 0.09 0.1 0.16 0.2 0.4 0.5 0.53 0.55 0.65 0.7 0.75 0.8 0.9 1
## x_vals   1.00 2.00 3.00 4.0 1.00 2.0 3.0 4.0 1.00 2.00 3.00 4.0 1.00 2.0 3.0 4
## y_vals   1.00 1.00 1.00 1.0 2.00 2.0 2.0 2.0 3.00 3.00 3.00 3.0 4.00 4.0 4.0 4

position <- min(which( u < cum_prob)) # compare u to cum_prob and finds the first value that is greater
position

## [1] 2
x_vals[position] # value to save for X

## [1] 2
y_vals[position] # value to save for Y

## [1] 1

```

```

# one value
u <- 0.89
print(rbind(cum_prob, x_vals, y_vals))

##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14] [,15] [,16]
## cum_prob 0.02 0.05 0.09 0.1 0.16 0.2 0.4 0.5 0.53 0.55 0.65 0.7 0.75 0.8 0.9 1
## x_vals   1.00 2.00 3.00 4.0 1.00 2.0 3.0 4.0 1.00 2.00 3.00 4.0 1.00 2.0 3.0 4
## y_vals   1.00 1.00 1.00 1.0 2.00 2.0 2.0 2.0 3.00 3.00 3.00 3.0 4.00 4.0 4.0 4

position <- min(which( u < cum_prob))
position

## [1] 15
x_vals[position] # value to save for X

## [1] 3
y_vals[position] # value to save for Y

## [1] 4

```

Code to generate many values

```
set.seed(1)
n <- 10 ^ 4
U <- matrix(runif(n)) # a matrix of U

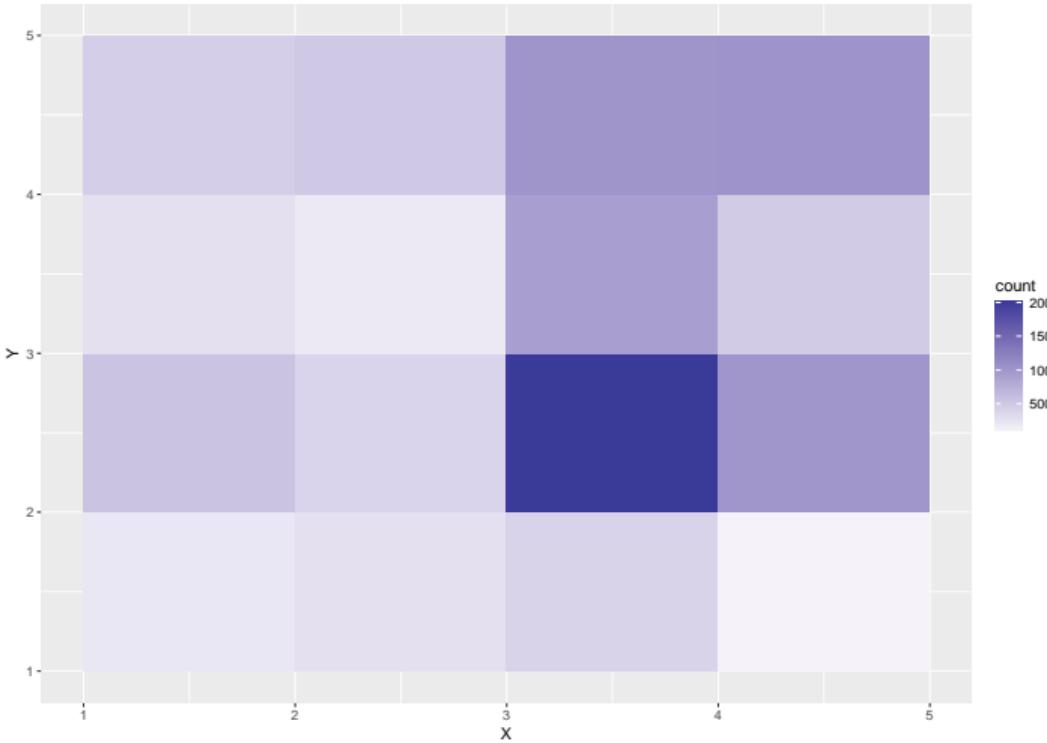
# a function to find the position based on u
position_func <- function(u){
  min(which( u < cum_prob))
}

positions <- apply(U, 1, position_func)

X <- x_vals[positions]
Y <- y_vals[positions]

df <- data.frame(X, Y, U)
```

```
library(ggplot2)
# for some reason binwidth = 1 produced only 3 bins
d <- ggplot(df, aes(X, Y)) + geom_bin2d(binwidth = 0.999) + scale_fill_gradient2()
d
```



Bivariate Normal Distribution

See Professor Tsiang's notes